Hochschule München University of Applied Sciences

Faculty of Computer Science and Mathematics

# parSAT: Parallel Solving of Floating-Point Satisfiability

**FROM 2025** 

Markus Krahl, Matthias Güdemann, and Stefan Wallentowitz



#### **Motivation**

- Embedded Systems control safety-critical processes by calculations of Floating Point (FP) numbers
- Most often FP-Implementations based on the IEEE 754 Standard are used
- FP numbers only approximate real values and do not fulfill mathematical rules (associativity, addition)
- Due to this behavior, FP algorithms might produce unexpected or incorrect results
- SMT solvers play a crucial role in verifying FP computations
- Current state-of-the-art SMT solvers convert FP arithmetic into propositional logic through bit- or word-blasting
- This overhead may result in limited capacity for reasoning about non-linear FP constraints
- Ideas for parSAT:
  - Exploit dedicated hardware support for FP computations (FP Units in CPUs)
  - Improve parallel solving of FP SMT equations



#### **Related Work**

- parSAT extends the concepts presented in XSat and goSAT:
  - XSat:
    - Foundations for translating FP constraints into Global Optimization (GO) problem
    - Emits the optimization function as C-code, compiles it, and it reloads it as Python-Module
    - Applies Basin Hopping (BH) from the Scipy-Package to find the global minimum
  - goSAT:
    - Based on XSat
    - Applies De-Morgan's Law when translating an SMT equation into a GO problem
    - Performs a Just-In-Time compilation of the optimization function
    - Employs various GO routines of the NLOpt library for finding the global minimum
- Research investigating portfolio-settings with SMT solvers



#### **Contributions**

- *parSAT*, an integrated tool that
  - represents a semi-decision procedure for SMT equations in FP theory
  - supports a broader scope of the SMTLIB2-standard
  - models the effect of infinite FP numbers in the generated optimization function
  - performs a portfolio-based minimization by using a multicore implementation
  - enables the integration of any GO algorithm
- Evaluation of parSAT with current state-of-the-art SMT solvers on different benchmarks



#### **Theoretical Background**

- Convert a quantifier-free FP SMT equation F into an optimization function G
  - If x corresponds to a satisfiable assignment of F than G(x) must be zero
  - Otherwise, **G** returns a positive distance value to the global minimum at zero
- Translation of negated comparison operators for FP numbers must be considered:

$$\neg (r_a < r_b)$$
  $(r_a \ge r_b)$   $\neg (f_a < f_b)$   $(f_a \ge f_b)$  where  $f_b = NaN$ 

• **F** in conjunctive normal form with removed negations (De-Morgan) is translated to **G**:

$$F_{CD}(\overrightarrow{x}) = \bigwedge_{i \in I} \bigvee_{j \in J} e_{i,j} \bowtie_{i,j,n} e'_{i,j}$$
 
$$G(\overrightarrow{x}) = \sum_{i \in I} \prod_{j \in J} d(\bowtie_{i,j,n}, e_{i,j}, e'_{i,j})$$

Subscript n indicates if a negation of comparison operator (bowtie) occurred



#### **Theoretical Background**

• Definitions for  $d(\bowtie_{i,j,n}, e_{i,j}, e'_{i,j})$ 

$$\begin{split} d(==_0,e_1,e_2) &= d(\neq_1,e_1,e_2) = \theta(e_1,e_2) \\ d(\neq_0,e_1,e_2) &= d(==_1,e_1,e_2) = e_1 \neq e_2 ? 0 : 1 \\ d(<_0,e_1,e_2) &= e_1 < e_2 ? 0 : \theta(e_1,e_2) + 1 \\ d(<_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(\geq_0,e_1,e_2) \\ d(\leq_0,e_1,e_2) &= e_1 \leq e_2 ? 0 : \theta(e_1,e_2) \\ d(\leq_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(>_0,e_1,e_2) \\ d(>_0,e_1,e_2) &= e_1 > e_2 ? 0 : \theta(e_1,e_2) + 1 \\ d(>_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(\leq_0,e_1,e_2) \\ d(\geq_0,e_1,e_2) &= e_1 \geq e_2 ? 0 : \theta(e_1,e_2) \\ d(\geq_0,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(<_0,e_1,e_2) \\ d(\geq_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(<_0,e_1,e_2) \\ d(\geq_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(<_0,e_1,e_2) \\ d(\geq_1,e_1,e_2) &= isnan(e_1) \lor isnan(e_2) ? 0 : d(<_0,e_1,e_2) \\ \end{pmatrix}$$

where


$$\theta(e_1, e_2) = isnan(e_1) \vee isnan(e_2) ? 1 : (e_1 == e_2 ? 0 : (|bits(e_1) - bits(e_2)|))$$




### **Example**

$$t(x) = -1(x+2)^2 - 2$$


**ExP(1)**: *x* ∈ *FP*  $\wedge$  *t*(*x*)  $\geq$  −2.











































#### **Incompleteness of parSAT**

- If parSAT finds a global minimum at zero, these coordinates represent a satisfiable solution to the initial SMT equation
- If *parSAT* reaches an evaluation limit (maximal number of calling the evaluation function exceeded) it is unknown whether the optimization function has a global minimum at zero (satisfiable) or not (unsatisfiable)
- The generated optimization function may not be smooth and therefore not differentiable
- Only derivative-free GO methods can be used
- Only when all possible inputs were evaluated and no solution was found, parSAT could safely determine UNSAT



## Distribution of fastest GO methods when evaluating the Griggio Benchmark

	BH	CRS2	ISRES
parsat $BH_1$ , $CRS2_1$ , $ISRES_1$	46.6%	27.9%	25.5%
parsat $BH_2$ , $CRS2_2$ , $ISRES_2$	59.1%	25.2%	15.7%
parsat BH <sub>3</sub> , CRS2 <sub>3</sub> , ISRES <sub>3</sub>	71.6%	17.3%	11.1%
parsat BH <sub>4</sub> , CRS2 <sub>4</sub> , ISRES <sub>4</sub>	74.0%	17.3%	8.7%
parsat BH <sub>5</sub> , CRS2 <sub>5</sub> , ISRES <sub>5</sub>	77.5%	15.5%	7.0%
parsat BH <sub>6</sub> , CRS2 <sub>6</sub> , ISRES <sub>6</sub>	83.6%	10.0%	6.4%
parsat BH <sub>7</sub> , CRS2 <sub>7</sub> , ISRES <sub>7</sub>	87.0%	7.7%	5.3%
parsat BH <sub>8</sub> , CRS2 <sub>8</sub> , ISRES <sub>8</sub>	87.5%	7.0%	5.5%
parsat BH <sub>9</sub> , CRS2 <sub>9</sub> , ISRES <sub>9</sub>	89.6%	5.1%	5.3%
parsat $BH_{10}$ , $CRS2_{10}$ , $ISRES_{10}$	91.2%	4.1%	4.7%




# Full comparison of parSAT with state-of-the-art SMT solvers (Griggio Benchmark)

	SAT	UNSAT	timeout / evaluation limit	errors	average SAT runtime
$parSAT_{best}BH_{14}, CRS2_3, ISRES_3$	102	0	112	0	0.1
bitwuzla	108	76	30	0	32.71
cvc5	104	76	30	4	34.01
MathSAT	100	69	45	0	25.86
Z3	75	56	54	29	88.7



# Solving time of parSAT and state-of-the-art SMT solvers for SAT equations (Griggio Benchmark)





#### Introduction to the 2019-Guedemann Benchmark

- Re-Implementation of natural logarithm (In') and exeponential (exp') function
- The 2019-Guedemann Benchmark includes SMT equations to verify that the following properties hold for the re-implemented algorithms

$$\forall x \in \mathbb{R}: \quad x > 0 \implies \exp'(\ln' x) \approx x$$

$$\forall x \in \mathbb{R}: \quad x > 0 \implies \ln'(\exp' x) \approx x$$

$$\forall x, y \in \mathbb{R}: \quad x, y \in [0, 1] \implies \exp'(x + y) \approx \exp'(x) \cdot \exp'(y)$$

$$\forall x \in \mathbb{R}: \quad x, y > 0 \implies \ln'(x \cdot y) \approx \ln'(x) \cdot \ln'(y)$$

$$\forall x \in \mathbb{R}: \quad x, y \in [0, 1] \implies img(x^y) = img(\exp'(y \cdot \ln'(x))) \approx [0, 1]$$



# Full comparison of parSAT with state-of-the-art SMT solvers (2019-Guedemann Benchmark)

	SAT	UNSAT	timeout / evaluation limit	errors	average SAT runtime
$parSAT_{best}BH_{14}, CRS2_3, ISRES_3$	10	0	3	0	0.1
bitwuzla	10	2	1	0	37.88
cvc5	4	1	8	0	180.09
MathSAT	4	1	8	0	22.8
Z3	0	0	4	9	-



# Potential effect of combining parSAT with bitwuzla (Griggio Benchmark)

	SAT	UNSAT	timeout / evaluation limit	errors	average SAT runtime
$parSAT_{best}BH_{14}, CRS2_3, ISRES_3$	102	0	112	0	0.1
bitwuzla	108	76	30	0	32.71
$parSAT_{best}$ + bitwuzla	113	76	25	0	12.75



#### **Conclusion and Outlook**

- parSAT presents a considerable complementary alternative to current state-of-the-art SMT solvers
- parSAT was able to handle a variety of SMT equations from commonly used, representative benchmarks
- The modular architecture of parSAT enables integration of further GO algorithms
- Due to the necessary evaluation limit for the optimization function, *parSAT* is a semi-decision procedure that underapproximates the search space and cannot prove *UNSAT*
- parSAT is limited to only find solutions consisting of finite FP values
- Extend parSAT to utilize other platforms for running concurrent GO methods such as GPUs or (cloud) computing clusters
- Enable information sharing between *parSAT* with SMT solvers in a complementary setting
- Investigate parSAT's potential in the problem domain of SMT model counting

